Nancy, the lazy web site builder
User's guide

Reuben Thomas

12th November 2008

1 Introduction

Nancy is a simple web site builder that glues together HTML and other fragments to make pages, and allows fragments to be specialised for particular pages. Fragments can be files or generated programmatically.

Nancy is a command-line tool.

2 Invocation

Nancy takes three (or optionally four) arguments:

```
nancy.pl SOURCE DESTINATION TEMPLATE [BRANCH]
```

where `SOURCE` is the source directory tree, `DESTINATION` is the directory to which the resulting HTML pages will be written, `TEMPLATE` is the name of the template file, and the optional `BRANCH` gives the sub-directory of the `SOURCE` tree to process (if it is omitted, the complete source tree are processed).

If you supply the optional flag `-list-files`, or `-l`, the files read by Nancy will be listed on standard error, and a message will be inserted into the output directly before the contents of the include giving the file name.

3 Operation

Nancy produces the finished pages according to the following algorithm:

1. For each leaf directory of the source tree, start from that directory.

2. Set the initial text to `$include{TEMPLATE}`.

3. Repeatedly scan the text for a command and replace it by its output, until no more commands are found.

4. Write the resultant text to a file: for each directory `SOURCE/LEAF_PATH` the output file is `DESTINATION/LEAF_PATH.html`.

The reason that only leaf directories correspond to pages is to ensure that every page can be specialised without affecting any other page. It is advisable to ensure that every non-leaf directory has an `index.html` sub-directory (or some other valid index page name), so that there are no "missing" URLs in the resulting site.

A command takes the form

`$COMMAND{ARGUMENT, ...}`

Nancy recognises these commands:

`$include{FILE}` Replace the command with the contents of the given file.

`$page{}` Replace the command with the URL of the page under construction, relative to the root of the site.

`$root{}` Replace the command with the relative URL to the root of the site from the page under construction. This means that every link in a site can be written relative to the current page, either explicitly (which is a good way to link to pages related to the current page, as such links do not need to be rewritten if the related pages are moved together within the site), or implicitly as `<a href="root{}/path/to/page.html">`. Hence the site's base URL can be changed without needing to change any intra-site links.

`$run{PROGRAM[, ARGUMENT, ...]}` Replace the command with the output of the given program. The command is run with any arguments given in the command.

Only one guarantee is made about the order in which commands are processed: if one command is nested inside another, the inner command will be processed first. (The order only matters for —`$run` commands; if you nest them, you have to deal with this potential pitfall.)

To find the file `FILE_PATH` specified by an `$include` command, Nancy proceeds thus:

1. Look in `SOURCE/BRANCH/LEAF_PATH/FILE_PATH`.

2. If the file is not found, remove the final directory from `LEAF_PATH` and try again, until `LEAF_PATH` is empty.

3. Finally, try looking in `SOURCE/BRANCH/FILE_PATH`.

For example, if `SOURCES` is `/dir`, `BRANCH` is omitted and `FILE_PATH` is `foo/bar/baz`, and Nancy is trying to find `file.html`, it will try the following directories, in order:

1. `/dir/foo/bar/baz/file.html`

2. `/dir/foo/bar/file.html`

3. `/dir/foo/file.html`

4. `/dir/file.html`

4 Example

Suppose a web site has the following page design, from top to bottom: logo, navigation menu, breadcrumb trail, page body.

Most of the elements are the same on each page, but the breadcrumb trail has to show the canonical path to each page, and the logo is bigger on the home page.

Suppose further that the web site has the following structure, where each line corresponds to a page:

- Home page

- People
    - Jo Bloggs
    - Hilary Pilary
    - ...

- Places
    - Vladivostok
    - Timbuktu
    - ...

The basic page template looks something like this:

```
<html>
  <link href="style.css" rel="stylesheet" type="text/css">
  <title>$include{title}</title>
  <body>
    <div class="logo">$include{logo.html}</div>
    <div class="menu">$include{menu.html}</div>
    <div class="breadcrumb">$include{breadcrumb.html}</div>
    <div class="main">$include{main.html}</div>
  </body>
</html>
```

Making the menu an included file is not strictly necessary, but, as in programming, makes the HTML fragments easier to read. The pages will be laid out as follows:

- /
    - index.html
    - people/
        * index.html
        * jo_bloggs.html
        * hilary_pilary.html
    - places/

- ∗ `index.html`
- ∗ `vladivostok.html`
- ∗ `timbuktu.html`

The corresponding source files will be laid out as follows. This may look a little confusing at first, but note the similarity to the HTML pages, and hold on for the explanation!

- `source/`
  - `template.html` (the template shown above)
  - `menu.html`
  - `logo.html`
  - `breadcrumb.html`
  - `index.html/`
    - ∗ `main.html`
    - ∗ `logo.html`
  - `people/`
    - ∗ `breadcrumb.html`
    - ∗ `index.html/`
      - · `main.html`
    - ∗ `jo_bloggs.html/`
      - · `main.html`
    - ∗ `hilary_pilary.html/`
      - · `main.html`
  - `places/`
    - ∗ `breadcrumb.html`
    - ∗ `index.html/`
      - · `main.html`
    - ∗ `vladivostok.html/`
      - · `main.html`
    - ∗ `timbuktu.html/`
      - · `main.html`

Note that there is only one menu fragment (the main menu is the same for every page), while each section has its own breadcrumb trail (`breadcrumb.html`), and each page has its own content (`main.html`).

To build the site, Nancy is invoked as:

```
nancy.pl source template.html dest
```

Now consider how Nancy builds the page whose URL is `vladivostok.html`. According to the rules given in Section 3, Nancy will look first for files in `source/places/vladivostok.html`, then in `source/places`, and finally in `source`. Hence, the actual list of files used to assemble the page is:

- `source/template.html`

4

- `source/logo.html`

- `source/menu.html`

- `source/places/breadcrumb.html`

- `source/places/vladivostok.html/main.html`

For the site's index page, the file `index.html/logo.html` will be used for the logo fragment, which can refer to the larger graphic desired.

This scheme, though simple, is surprisingly flexible; this simple example has covered all the standard techniques for Nancy's use.